

# **Automated Resume Screening Using NLP: Evaluating Fairness and Mitigating Bias**

**A PROJECT REPORT**

*Submitted by*

**Yeshika (23BDA70132)**

**Bedant Singhania (23BDA70122)**

**Gaurav Sharma (23BDA70050)**

**Javed Akhtar (23BDA70114)**

**Manshvi Sharma (23BDA70088)**

*in partial fulfillment for the award of the degree of*

**BACHELOR OF ENGINEERING**

**IN**

**COMPUTER SCIENCE AND ENGINEERING**



**CHANDIGARH UNIVERSITY  
MOHALI, PUNJAB**

**APRIL 2026**



## BONAFIDE CERTIFICATE

Certified that this project report “**Automated Resume Screening Using NLP: Evaluating Fairness and Mitigating Bias**” is the bonafide work of **Yeshika (23BDA70132), Bedant Singhania (23BDA70122), Gaurav Sharma (23BDA70050), Javed Akhtar (23BDA70114), and Manshvi Sharma (23BDA70088)** who carried out the project work under our supervision.

**SIGNATURE**

**SIGNATURE**

**Name of the Head of the Department**

**Amanpreet Kaur**

**HEAD OF THE DEPARTMENT**

**SUPERVISOR**

Department of Computer Science and Engineering

Assistant Professor

Chandigarh University, Mohali

Department of Computer Science and Engineering

Submitted for the project viva-voce examination held on \_\_\_\_\_

**INTERNAL EXAMINER**

**EXTERNAL EXAMINER**

## ACKNOWLEDGEMENT

We would like to express our sincere and heartfelt gratitude to our project supervisor, **Amanpreet Kaur**, Assistant Professor, Department of Computer Science and Engineering, Chandigarh University, for her valuable guidance, continuous encouragement, and constructive suggestions throughout the development of this project titled “*Automated Resume Screening Using NLP: Evaluating Fairness and Mitigating Bias.*”

Her guidance helped us understand the theoretical and practical aspects of Natural Language Processing, machine learning models, fairness evaluation, ethical artificial intelligence, and bias mitigation techniques. Her constant support motivated us to improve the quality of our research work and present it in a structured and professional manner.

We are highly thankful to the **Department of Computer Science and Engineering, Chandigarh University**, for providing us with the required academic environment, learning resources, laboratory support, and technical exposure needed to complete this project successfully.

We would also like to thank all faculty members of the department for their academic support and encouragement. Their lectures, discussions, and technical guidance helped us strengthen our understanding of machine learning, deep learning, data analysis, and research methodology.

We are grateful to our classmates and friends for their cooperation, feedback, and useful suggestions during different stages of the project. Their support helped us improve the overall clarity and quality of our work.

Finally, we express our deepest gratitude to our parents and family members for their continuous motivation, patience, and moral support throughout the completion of this project.

**Yeshika**

**Bedant Singhania**

**Gaurav Sharma**

**Javed Akhtar**

**Manshvi Sharma**

# TABLE OF CONTENTS

|          |                                                                        |           |
|----------|------------------------------------------------------------------------|-----------|
| <b>1</b> | <b>Introduction</b>                                                    | <b>1</b>  |
| 1.1      | Background of the Study . . . . .                                      | 1         |
| 1.2      | Motivation of the Study . . . . .                                      | 1         |
| 1.3      | Research Gap . . . . .                                                 | 2         |
| 1.4      | Problem Statement . . . . .                                            | 4         |
| 1.5      | Objectives of the Study . . . . .                                      | 4         |
| 1.6      | Scope of the Study . . . . .                                           | 5         |
| 1.7      | Methodology Overview . . . . .                                         | 6         |
| 1.8      | Proposed System Architecture . . . . .                                 | 7         |
| 1.9      | Project Timeline . . . . .                                             | 7         |
| 1.10     | Contribution of the Study . . . . .                                    | 7         |
| 1.11     | Applications of the Proposed System . . . . .                          | 8         |
| 1.12     | Significance of the Study . . . . .                                    | 9         |
| 1.13     | Organization of the Report . . . . .                                   | 9         |
| <b>2</b> | <b>Literature Survey</b>                                               | <b>10</b> |
| 2.1      | Introduction . . . . .                                                 | 10        |
| 2.2      | Evolution of Resume Screening Systems . . . . .                        | 10        |
| 2.3      | Traditional Machine Learning Approaches for Resume Screening . . . . . | 11        |
| 2.4      | Transformer-Based Resume Screening Models . . . . .                    | 11        |
| 2.5      | Bias in Automated Recruitment Systems . . . . .                        | 12        |
| 2.6      | Fairness Metrics in Machine Learning . . . . .                         | 12        |
| 2.7      | Bias Mitigation Techniques . . . . .                                   | 13        |
| 2.8      | Dual-Track NLP Resume Screening Pipelines . . . . .                    | 14        |
| 2.9      | Comparison Between Classical and Transformer Models . . . . .          | 14        |
| 2.10     | Research Gap . . . . .                                                 | 15        |
|          | <b>CHAPTER 3: DESIGN FLOW / PROCESS</b>                                | <b>16</b> |

|          |                                                                    |           |
|----------|--------------------------------------------------------------------|-----------|
| 3.1      | Evaluation and Selection of Specifications / Features . . . . .    | 16        |
| 3.2      | Analysis and Feature Finalization Subject to Constraints . . . . . | 17        |
| 3.3      | Model Training Strategy . . . . .                                  | 21        |
| 3.4      | Transformer Fine-Tuning Configuration . . . . .                    | 22        |
| 3.5      | Bias Detection Methodology . . . . .                               | 23        |
| 3.6      | Bias Mitigation Strategy . . . . .                                 | 23        |
| 3.7      | Implementation Workflow . . . . .                                  | 24        |
| 3.8      | System Scalability Considerations . . . . .                        | 24        |
| <b>4</b> | <b>Results Analysis and Validation</b>                             | <b>26</b> |
| 4.1      | Model Performance . . . . .                                        | 26        |
| 4.2      | Bias Mitigation Results . . . . .                                  | 29        |
| 4.3      | Validation . . . . .                                               | 30        |
| 4.4      | Interpretation . . . . .                                           | 32        |
| <b>5</b> | <b>Conclusion and Future Work</b>                                  | <b>35</b> |
| 5.1      | Conclusion . . . . .                                               | 35        |
| 5.2      | Validation of Project Objectives . . . . .                         | 36        |
| 5.3      | Research Contributions of the Project . . . . .                    | 37        |
| 5.4      | Practical Significance of the Proposed Framework . . . . .         | 38        |
| 5.5      | Ethical Impact and Responsible AI Contribution . . . . .           | 38        |
| 5.6      | Deviation from Expected Results . . . . .                          | 39        |
| 5.7      | Limitations of the Study . . . . .                                 | 39        |
| 5.8      | Future Work . . . . .                                              | 40        |
| 5.8.1    | Testing on Real-World Resume Datasets . . . . .                    | 40        |
| 5.8.2    | Explainable AI Dashboard Development . . . . .                     | 40        |
| 5.8.3    | Intersectional Bias Analysis . . . . .                             | 40        |
| 5.8.4    | Web-Based Resume Screening Platform . . . . .                      | 40        |
| 5.8.5    | Multilingual Resume Processing . . . . .                           | 40        |
| 5.8.6    | Human-in-the-Loop Decision Support . . . . .                       | 40        |
| 5.9      | Scalability of the Proposed Framework . . . . .                    | 40        |
| 5.10     | Final Summary . . . . .                                            | 41        |
| <b>A</b> | <b>Appendix</b>                                                    | <b>43</b> |

|       |                                               |    |
|-------|-----------------------------------------------|----|
| A.1   | Sample Resume Fields . . . . .                | 43 |
| A.2   | Dataset Preparation Workflow . . . . .        | 43 |
| A.3   | User Manual . . . . .                         | 44 |
| A.3.1 | System Requirements . . . . .                 | 44 |
| A.3.2 | Steps to Run the Project . . . . .            | 45 |
| A.3.3 | Execution Workflow of the Framework . . . . . | 45 |
| A.3.4 | System Outputs . . . . .                      | 46 |
| A.4   | Fairness Evaluation Module . . . . .          | 46 |
| A.5   | Bias Mitigation Module . . . . .              | 46 |
| A.6   | Achievements . . . . .                        | 47 |

# LIST OF FIGURES

|     |                                                                          |    |
|-----|--------------------------------------------------------------------------|----|
| 1   | Graphical Abstract of the Proposed NLP Resume Screening System . . . . . |    |
| 1.1 | Architecture of Automated Resume Screening System . . . . .              | 8  |
| 1.2 | Project Development Timeline . . . . .                                   | 9  |
| 2.1 | Performance Comparison of Resume Classification Models . . . . .         | 11 |
| 2.2 | BERT / RoBERTa Fine-Tuning Pipeline for Resume Classification . . . . .  | 12 |
| 2.3 | DIR Before and After Bias Mitigation . . . . .                           | 13 |
| 2.4 | Dual-Track NLP Resume Screening Pipeline . . . . .                       | 14 |
| 4.1 | Accuracy Comparison of Resume Screening Models . . . . .                 | 27 |
| 4.2 | DIR Before and After Bias Mitigation . . . . .                           | 30 |
| 4.3 | Validation Workflow of the Resume Screening Framework . . . . .          | 32 |
| 4.4 | Hybrid Interpretation Strategy for Resume Classification . . . . .       | 33 |

## LIST OF TABLES

|     |                                                                                          |    |
|-----|------------------------------------------------------------------------------------------|----|
| 1.1 | Comparison Between Existing Resume Screening Approaches and Proposed Framework . . . . . | 3  |
| 1.2 | Methodology Pipeline of the Proposed Resume Screening Framework . . . . .                | 7  |
| 2.1 | Comparison of Resume Classification Models . . . . .                                     | 15 |
| 3.1 | Feature Evaluation Matrix for Resume Screening Framework . . . . .                       | 16 |
| 3.2 | Feature Finalization after Constraint Analysis . . . . .                                 | 20 |
| 3.3 | Constraint-Based Feature Selection Considerations . . . . .                              | 21 |
| 3.4 | Implementation Workflow of Resume Screening Framework . . . . .                          | 24 |
| 4.1 | Model Performance Comparison . . . . .                                                   | 26 |
| 4.2 | Fairness Metrics Before Mitigation . . . . .                                             | 28 |
| 4.3 | Fairness Metrics After Mitigation . . . . .                                              | 29 |
| 5.1 | Validation of Project Objectives . . . . .                                               | 37 |



# ABSTRACT

As the number of software engineering job applications increases rapidly, organizations are increasingly adopting automated resume screening systems to manage candidate pipelines efficiently and reduce manual recruitment workload. Traditional keyword-based filtering approaches are limited in their ability to interpret contextual relationships within resumes, often leading to inaccurate candidate shortlisting decisions. Recent advances in Natural Language Processing (NLP) have significantly improved automated screening capabilities by enabling machines to understand resume content semantically, extract relevant skills and qualifications, compare candidate profiles with job requirements, and support faster and more consistent recruitment decision-making processes.

However, despite these advantages, automated screening systems may unintentionally reproduce historical bias present in training datasets. Bias can appear through gendered language patterns, culturally associated names, age indicators, educational background signals, and other proxy variables embedded in resume text. Such hidden correlations may result in unfair rejection of qualified candidates from underrepresented demographic groups, raising important ethical and regulatory concerns regarding the deployment of Artificial Intelligence in recruitment environments.

This project presents a Natural Language Processing-based automated resume screening system designed specifically for evaluating candidate suitability for software engineering roles while simultaneously measuring and mitigating demographic bias. The proposed framework integrates classical machine learning classifiers including Logistic Regression, Support Vector Machine, and Random Forest using TF-IDF feature representations, together with transformer-based deep learning architectures such as BERT and RoBERTa to capture contextual semantic relationships within resume content. A synthetic dataset consisting of 1,200 structured resumes

was developed to enable controlled experimentation across multiple protected demographic attributes including gender, age group, race or ethnicity, and name-origin indicators.


Fairness evaluation was performed using widely accepted quantitative metrics including Disparate Impact Ratio (DIR), Demographic Parity Difference (DPD), and Equalized Odds Difference (EOD), allowing systematic identification of demographic imbalance in classification outcomes. Experimental results demonstrate that transformer-based architectures achieve significantly higher prediction accuracy compared to classical machine learning approaches, with RoBERTa achieving the highest classification accuracy of 93.2%. However, these models also exhibited stronger sensitivity to demographic proxy variables prior to mitigation.

To address this limitation, bias mitigation techniques including sample re-weighting, adversarial debiasing, and counterfactual data augmentation were integrated into the screening pipeline. After mitigation, Disparate Impact Ratio values increased beyond recommended fairness thresholds across gender, age, race or ethnicity, and name-based dimensions, demonstrating the effectiveness of the proposed framework in achieving balanced and responsible automated screening outcomes.

Overall, the results confirm that combining contextual transformer-based resume understanding with fairness-aware machine learning techniques enables the development of accurate, interpretable, and ethically reliable recruitment support systems. The proposed framework therefore contributes toward the advancement of trustworthy AI-assisted hiring technologies suitable for deployment in intelligent recruitment analytics environments.

**Keywords:** Natural Language Processing, Resume Screening, Bias Detection, Fairness, BERT, RoBERTa, TF-IDF, Disparate Impact Ratio, Ethical AI.

# GRAPHICAL ABSTRACT



**Figure 1: Graphical Abstract of the Proposed NLP Resume Screening System**

## ABBREVIATIONS

| Abbreviation | Full Form                                               |
|--------------|---------------------------------------------------------|
| AI           | Artificial Intelligence                                 |
| NLP          | Natural Language Processing                             |
| ML           | Machine Learning                                        |
| DL           | Deep Learning                                           |
| TF-IDF       | Term Frequency–Inverse Document Frequency               |
| SVM          | Support Vector Machine                                  |
| LR           | Logistic Regression                                     |
| RF           | Random Forest                                           |
| BERT         | Bidirectional Encoder Representations from Transformers |
| RoBERTa      | Robustly Optimized BERT Pretraining Approach            |
| DIR          | Disparate Impact Ratio                                  |
| DPD          | Demographic Parity Difference                           |
| EOD          | Equalized Odds Difference                               |
| CDA          | Counterfactual Data Augmentation                        |
| NER          | Named Entity Recognition                                |
| EEOC         | Equal Employment Opportunity Commission                 |
| EU           | European Union                                          |
| HR           | Human Resources                                         |

# SYMBOLS

| Symbol             | Meaning                                                  |
|--------------------|----------------------------------------------------------|
| $\hat{Y}$          | Predicted output label                                   |
| $Y$                | Actual output label                                      |
| $A$                | Protected attribute                                      |
| $P(\hat{Y} = 1 A)$ | Probability of positive prediction for a protected group |
| $TPR$              | True Positive Rate                                       |
| $FPR$              | False Positive Rate                                      |
| $W_i$              | Weight of sample $i$                                     |
| $\theta$           | Model parameters                                         |
| $\mathcal{L}$      | Loss function                                            |
| $\lambda$          | Regularization parameter                                 |
| $n$                | Number of samples                                        |
| $d$                | Number of features                                       |

# **CHAPTER-1**

## **Introduction**

### **1.1 Background of the Study**

Recruitment is a critical process in modern organizations, especially in technical domains such as software engineering where large volumes of applications are received for a limited number of positions. Traditional resume screening methods rely heavily on manual evaluation by recruiters, which is time-consuming, inconsistent, and prone to human bias.

Recent advancements in Artificial Intelligence (AI) and Natural Language Processing (NLP) have enabled automated resume screening systems capable of analyzing candidate resumes efficiently. These systems extract relevant features such as skills, qualifications, experience, and certifications from resumes and classify candidates according to job requirements.

However, automated hiring systems trained on historical recruitment datasets may unintentionally inherit demographic biases related to gender, ethnicity, age, or candidate names. Such biases reduce fairness and transparency in recruitment decisions. Therefore, designing intelligent resume screening systems that ensure both high prediction accuracy and fairness is essential.

This research proposes an automated resume screening framework using TF-IDF, BERT, and RoBERTa models along with bias detection and mitigation techniques.

### **1.2 Motivation of the Study**

Organizations receive thousands of resumes for a single job posting, making manual screening inefficient and error-prone. Although automated screening systems improve efficiency, many existing systems fail to ensure fairness in decision-making.

The motivation behind this research is to develop a fair and intelligent resume screening

system that improves classification accuracy while minimizing demographic bias in recruitment decisions.

### 1.3 Research Gap

Existing resume screening systems primarily focus on improving classification accuracy without addressing fairness issues. Transformer-based NLP models such as BERT and RoBERTa provide high accuracy but may amplify bias present in training datasets. As a result, automated recruitment systems trained on historical hiring data can unintentionally reproduce demographic inequalities related to gender, age group, race or ethnicity, and name-origin indicators.

Most traditional resume screening approaches rely on keyword matching or TF-IDF feature representations, which are limited in capturing contextual semantic relationships within resume content. Although recent transformer-based models improve classification performance significantly, they often operate as black-box systems with limited interpretability and insufficient fairness monitoring mechanisms.

Another important limitation observed in existing literature is that fairness evaluation is frequently treated as a secondary validation step rather than being integrated directly into the model development pipeline. Many prior studies report classification accuracy improvements without systematically analysing demographic bias or applying mitigation strategies before deployment. This creates potential risks when such systems are used in real-world recruitment environments where ethical and regulatory compliance is critical.

Furthermore, current research often evaluates bias using only a single protected attribute, whereas practical hiring systems must consider multiple demographic dimensions simultaneously. There is also limited work available on combining classical machine learning models with transformer-based architectures within a single fairness-aware evaluation framework.

Therefore, there is a clear need for an integrated resume screening framework that:

- combines classical machine learning and transformer-based classification models,
- evaluates demographic bias using multiple fairness metrics,
- applies mitigation strategies during the prediction pipeline,
- maintains a balance between predictive accuracy and ethical reliability,

- supports scalable deployment in intelligent recruitment analytics systems.

To address these limitations, the present study proposes a hybrid Natural Language Processing–based resume screening framework that jointly performs candidate classification, bias detection, and fairness-aware optimization within a unified architecture.

**Table 1.1: Comparison Between Existing Resume Screening Approaches and Proposed Framework**

| Feature                  | Existing Systems                | Proposed Framework                       |
|--------------------------|---------------------------------|------------------------------------------|
| Primary Objective        | Accuracy-focused classification | Accuracy + fairness-aware classification |
| Feature Extraction       | Keyword matching / TF-IDF       | TF-IDF + Transformer embeddings          |
| Context Understanding    | Limited                         | Context-aware semantic modeling          |
| Fairness Evaluation      | Rarely included                 | Integrated using DIR, DPD, EOD           |
| Bias Mitigation          | Not commonly implemented        | Reweighting, CDA, adversarial debiasing  |
| Interpretability Support | Partial                         | Hybrid interpretable architecture        |
| Deployment Readiness     | Accuracy-oriented only          | Ethical AI–ready recruitment framework   |

The comparison presented above highlights that existing resume screening systems generally prioritize predictive performance without incorporating fairness-aware optimization strategies. The proposed framework addresses this gap by integrating classification accuracy, demographic bias detection, and mitigation techniques within a unified automated recruitment pipeline, thereby contributing toward the development of responsible and trustworthy AI-assisted hiring systems.



## 1.4 Problem Statement

Manual resume screening is inefficient for handling large datasets and often leads to inconsistent hiring decisions. Existing automated systems improve efficiency but may introduce bias against certain demographic groups.

Therefore, there is a need to develop an automated resume screening system capable of accurate classification while ensuring fairness across demographic attributes such as gender, ethnicity, and age.

## 1.5 Objectives of the Study

The primary objective of this study is to design and evaluate a fairness-aware automated resume screening framework using Natural Language Processing techniques for software engineering recruitment tasks. The system aims to improve classification accuracy while ensuring ethical reliability and demographic neutrality across candidate groups.

The specific objectives of the study are:

- To develop an NLP-based automated resume screening system
- To classify resumes based on job-role relevance
- To implement TF-IDF, BERT, and RoBERTa classification models
- To evaluate classification performance using standard metrics
- To detect demographic bias using fairness evaluation metrics
- To apply bias mitigation techniques
- To improve fairness without reducing prediction accuracy

In addition to the above objectives, the study also aims to:

- compare contextual transformer-based models with interpretable classical machine learning approaches,
- investigate the relationship between classification accuracy and fairness performance,
- evaluate demographic neutrality across multiple protected attributes simultaneously,

- design a hybrid architecture suitable for responsible AI-assisted recruitment systems,
- establish a scalable framework that can support future Applicant Tracking System integration.

## 1.6 Scope of the Study

The proposed system focuses on resume classification for software engineering job roles using Natural Language Processing techniques. The framework processes resume text data using both classical machine learning and transformer-based contextual embedding models in order to evaluate candidate suitability for technical positions.

Bias evaluation is performed across gender, ethnicity, and age attributes. These attributes were selected because they represent commonly studied demographic dimensions in fairness-aware machine learning research and are known to influence historical recruitment datasets.

The scope of the study includes:

- preprocessing unstructured resume text using NLP techniques,
- extracting semantic and statistical features using TF-IDF and transformer embeddings,
- training classification models using Logistic Regression, Support Vector Machine, Random Forest, BERT, and RoBERTa,
- evaluating prediction performance using accuracy, precision, recall, and F1-score metrics,
- analysing demographic fairness using DIR, DPD, and EOD indicators,
- applying bias mitigation strategies to improve fairness outcomes.

The system is limited to offline dataset evaluation and does not include real-time enterprise deployment. Additionally, multilingual resume processing and large-scale cloud-based integration are outside the scope of the current implementation and are considered future extensions of this work.

Despite these limitations, the framework provides a strong experimental foundation for fairness-aware automated recruitment analytics systems.

## 1.7 Methodology Overview

The methodology of the proposed system consists of multiple stages including dataset preparation, preprocessing, feature extraction, classification, fairness evaluation, and bias mitigation. These stages collectively ensure that the system achieves both high prediction accuracy and demographic neutrality across candidate groups.

Initially, resumes are cleaned and normalized using NLP preprocessing techniques such as tokenization and stop-word removal. Additional preprocessing operations include lower-casing, punctuation removal, and lemmatization to improve feature consistency across resume documents.

Feature extraction is performed using TF-IDF and transformer embeddings. TF-IDF captures statistical importance of keywords within resumes, while transformer-based representations capture contextual semantic relationships between candidate skills, experience descriptions, and qualifications.

Classification models are trained using machine learning and transformer-based architectures. Classical classifiers provide interpretable baseline performance, whereas transformer models improve semantic understanding and prediction robustness.

Fairness metrics are applied to detect bias, followed by mitigation techniques such as reweighting and counterfactual augmentation. These techniques adjust dataset representation and feature learning behaviour in order to reduce demographic imbalance without degrading classification performance.

The methodology pipeline followed in this study is summarized in Table 1.2.

**Table 1.2: Methodology Pipeline of the Proposed Resume Screening Framework**

| Stage                  | Description                                                             |
|------------------------|-------------------------------------------------------------------------|
| Dataset Preparation    | Collection and structuring of resume data for supervised classification |
| Preprocessing          | Tokenization, stop-word removal, normalization, and lemmatization       |
| Feature Extraction     | TF-IDF vectors and transformer-based contextual embeddings              |
| Classification         | LR, SVM, RF, BERT, and RoBERTa model training                           |
| Performance Evaluation | Accuracy, Precision, Recall, and F1-score computation                   |
| Fairness Evaluation    | DIR, DPD, and EOD metric calculation                                    |
| Bias Mitigation        | Reweighting, adversarial debiasing, and counterfactual augmentation     |
| Final Analysis         | Comparison of performance before and after mitigation                   |

The structured methodology ensures that classification accuracy improvements are achieved alongside fairness-aware optimization, thereby supporting the development of reliable and ethically responsible automated resume screening systems.

## 1.8 Proposed System Architecture


Figure 1.1 illustrates the architecture of the proposed automated resume screening framework.

## 1.9 Project Timeline

The timeline of project development is illustrated in Figure 1.2.

## 1.10 Contribution of the Study

- Development of an automated resume screening system using NLP



**Figure 1.1: Architecture of Automated Resume Screening System**

- Implementation of transformer-based classification models
- Detection of demographic bias using fairness metrics
- Application of mitigation techniques for improving fairness
- Comparative analysis between traditional ML and transformer models

## 1.11 Applications of the Proposed System

The proposed system can be applied in:

- Automated recruitment platforms
- HR analytics systems
- Resume ranking engines

| Phase   | Activity               |
|---------|------------------------|
| Phase 1 | Literature Review      |
| Phase 2 | Dataset Collection     |
| Phase 3 | Data Preprocessing     |
| Phase 4 | Feature Extraction     |
| Phase 5 | Model Training         |
| Phase 6 | Bias Detection         |
| Phase 7 | Bias Mitigation        |
| Phase 8 | Evaluation and Testing |
| Phase 9 | Documentation          |

**Figure 1.2: Project Development Timeline**

- Talent acquisition tools
- AI-based hiring assistants

## **1.12 Significance of the Study**

This study contributes toward developing ethical AI recruitment systems by addressing fairness challenges in automated hiring. The system improves classification accuracy while ensuring transparency and unbiased candidate evaluation.

## **1.13 Organization of the Report**

- Chapter 1 introduces the research background and objectives
- Chapter 2 presents literature review
- Chapter 3 explains methodology and implementation
- Chapter 4 discusses results and evaluation
- Chapter 5 concludes the study and future work

## **CHAPTER-2**

### **Literature Survey**

#### **2.1 Introduction**

The literature survey explores previous research in automated recruitment systems, Natural Language Processing (NLP) based resume screening, fairness-aware machine learning, bias detection techniques, and bias mitigation strategies. It provides the theoretical foundation for the proposed system and highlights limitations in existing approaches that motivate the present study.

Recent developments in transformer-based architectures such as BERT and RoBERTa have significantly improved text classification performance. However, concerns regarding fairness and demographic bias remain an active research challenge in AI-driven hiring systems.

#### **2.2 Evolution of Resume Screening Systems**

Early resume screening systems relied primarily on keyword matching techniques. These systems compared job descriptions with resumes using simple lexical similarity measures. Although computationally efficient, keyword-based systems failed to capture semantic meaning and contextual relationships between skills and experiences.

Later approaches introduced machine learning algorithms such as Logistic Regression, Support Vector Machines (SVM), and Random Forest classifiers using TF-IDF feature representations. These methods improved classification accuracy but still lacked contextual understanding of language.


Recent advancements in deep learning introduced transformer-based architectures such as BERT and RoBERTa, which significantly improved semantic understanding of resume content through contextual embeddings.

## 2.3 Traditional Machine Learning Approaches for Resume Screening

Several studies have applied classical machine learning algorithms for automated resume classification. Popular approaches include:

- Logistic Regression with TF-IDF features
- Support Vector Machines (SVM)
- Random Forest classifiers

These models perform well on structured datasets but struggle to capture contextual relationships between words. As shown in Figure 2.1, transformer-based models outperform traditional approaches in classification accuracy and fairness-aware evaluation.




**Figure 2.1: Performance Comparison of Resume Classification Models**

## 2.4 Transformer-Based Resume Screening Models

Transformer-based architectures such as BERT and RoBERTa have become the state-of-the-art methods for text classification tasks. These models use self-attention mechanisms to capture contextual relationships between words across long sequences.

Fine-tuning transformer models on resume datasets improves classification accuracy significantly compared to traditional machine learning methods. The architecture used for resume classification in this study is illustrated in Figure 2.2.





**Figure 2.2: BERT / RoBERTa Fine-Tuning Pipeline for Resume Classification**

## 2.5 Bias in Automated Recruitment Systems

Although automated hiring systems improve efficiency, they may unintentionally inherit bias from historical recruitment data. Studies have shown that machine learning models trained on biased datasets may discriminate against candidates based on gender, ethnicity, age, or even candidate names.

Such bias can negatively impact diversity in recruitment processes and reduce trust in AI-based hiring systems. Therefore, fairness-aware evaluation has become an essential requirement in modern recruitment frameworks.

## 2.6 Fairness Metrics in Machine Learning

Fairness metrics are used to evaluate whether machine learning models treat demographic groups equally. Common fairness metrics include:

- Disparate Impact Ratio (DIR)
- Demographic Parity Difference (DPD)
- Equal Opportunity Difference (EOD)

DIR measures whether the selection rate for protected groups differs significantly from privileged groups. According to fairness standards, a DIR value above 0.80 indicates acceptable fairness levels.

Figure 2.3 illustrates improvements in fairness metrics after applying mitigation techniques.




Figure 2.3: DIR Before and After Bias Mitigation

## 2.7 Bias Mitigation Techniques

Bias mitigation techniques are categorized into three main approaches:

1. Pre-processing methods
2. In-processing methods
3. Post-processing methods

Pre-processing methods modify training data before model training. In-processing methods modify learning algorithms during training. Post-processing methods adjust predictions after classification.

Common mitigation techniques include:

- Sample re-weighting
- Counterfactual data augmentation
- Adversarial debiasing


- Threshold optimization

These methods help improve fairness without significantly affecting classification accuracy.

## 2.8 Dual-Track NLP Resume Screening Pipelines

Modern resume screening frameworks combine classical machine learning pipelines with transformer-based architectures to improve performance and fairness evaluation simultaneously.

The dual-track pipeline used in this study integrates TF-IDF-based classification and transformer-based embeddings with fairness evaluation modules, as illustrated in Figure 2.4.



**Figure 2.4: Dual-Track NLP Resume Screening Pipeline**

## 2.9 Comparison Between Classical and Transformer Models

Table 2.1 presents a comparison between traditional machine learning models and transformer-based architectures used for resume screening tasks.

| Model Type          | Advantages                     | Limitations                      |
|---------------------|--------------------------------|----------------------------------|
| Logistic Regression | Fast training                  | Limited contextual understanding |
| SVM                 | Effective in high dimensions   | Sensitive to parameter tuning    |
| Random Forest       | Handles non-linear data        | Large memory requirement         |
| BERT                | Context-aware embeddings       | High computational cost          |
| RoBERTa             | Improved performance over BERT | Requires large datasets          |

**Table 2.1: Comparison of Resume Classification Models**

## 2.10 Research Gap

Most existing resume screening systems focus primarily on improving classification accuracy without addressing fairness issues. Traditional machine learning approaches lack contextual understanding, while transformer-based approaches may amplify bias present in training datasets.

There is a need for an integrated framework that combines transformer-based classification with fairness-aware evaluation and mitigation strategies across multiple demographic attributes such as gender, ethnicity, age, and candidate names.

The proposed study addresses this gap by developing a hybrid resume screening system that improves prediction performance while ensuring fairness in automated hiring decisions.

# CHAPTER 3

## DESIGN FLOW / PROCESS

### 3.1 Evaluation and Selection of Specifications / Features

The proposed automated resume screening framework integrates Natural Language Processing techniques, machine learning classification algorithms, transformer-based contextual embeddings, and fairness-aware optimization mechanisms. The system was designed after evaluating computational feasibility, classification performance, ethical fairness requirements, and scalability for future deployment in recruitment automation environments.

The selected feature set ensures that the screening system performs accurate candidate classification while maintaining demographic neutrality across gender, age groups, and name-based identity categories.

**Table 3.1: Feature Evaluation Matrix for Resume Screening Framework**

| Feature                     | Feasibility | Impact | Selected     |
|-----------------------------|-------------|--------|--------------|
| Resume Dataset Structuring  | High        | High   | Yes          |
| Text Preprocessing Pipeline | High        | High   | Yes          |
| TF-IDF Feature Extraction   | High        | Medium | Yes          |
| Transformer Embeddings      | Medium      | High   | Yes          |
| Classification Models       | High        | High   | Yes          |
| Fairness Evaluation Metrics | High        | High   | Yes          |
| Bias Mitigation Algorithms  | Medium      | High   | Yes          |
| Explainability Dashboard    | Medium      | Medium | Future Scope |
| ATS Integration             | Medium      | High   | Future Scope |

In addition to feasibility and performance considerations, the selected feature set was evaluated with respect to interpretability, fairness transparency, and deployment readiness. Classical machine learning models based on TF-IDF were retained to provide baseline comparison and explainable decision support, while transformer-based architectures such as BERT and RoBERTa were incorporated to capture contextual semantic relationships within resume text. The inclusion of fairness evaluation metrics and bias mitigation techniques ensures that the

screening framework aligns with ethical AI guidelines and supports unbiased candidate selection. Features requiring large-scale infrastructure, such as real-time ATS integration and explainable visualization dashboards, were identified as future extensions to maintain implementation feasibility within the scope of the current study.

### **3.2 Analysis and Feature Finalization Subject to Constraints**

The selection of system features was carried out after carefully analyzing technical feasibility, computational complexity, ethical fairness requirements, legal compliance considerations, and scalability constraints. Since automated resume screening systems directly influence recruitment decisions, it was essential to ensure that the final framework achieves a balance between predictive accuracy, model interpretability, fairness transparency, and deployment practicality. The objective was not only to develop a high-performance classification system but also to ensure that the screening decisions remain explainable, reliable, and ethically acceptable in real-world hiring environments.

During the feature selection process, priority was given to techniques that support structured resume interpretation while maintaining compatibility with both classical machine learning pipelines and transformer-based deep learning architectures. This hybrid selection strategy ensures that the proposed framework remains suitable for academic experimentation as well as scalable enterprise-level deployment in automated recruitment platforms.

#### **Selection of Classical Machine Learning Models**

Classical machine learning techniques such as TF-IDF with Logistic Regression, Support Vector Machine, and Random Forest were retained in the architecture to provide interpretable baseline predictions. These models enable transparent analysis of keyword-level importance and allow recruiters to understand how classification decisions are influenced by candidate skills, education, and experience descriptions.

The inclusion of classical models provides several advantages:

- They establish a baseline reference for comparing transformer-based model performance.
- They require lower computational resources during training and inference.
- They provide interpretable feature-weight relationships useful for recruiter validation.
- They improve system reliability when transformer inference resources are limited.

- They support faster experimentation during model evaluation stages.

Because interpretability remains a critical requirement in recruitment decision-making, these models play an important role in maintaining transparency within the screening pipeline.

### **Selection of Transformer-Based Architectures**

Transformer-based architectures including BERT and RoBERTa were selected to improve semantic understanding of resume content. Unlike TF-IDF representations, transformer models capture contextual relationships between words and phrases using attention-based language modeling mechanisms.

These architectures enable the framework to:

- understand variations in skill descriptions across resumes,
- capture contextual dependencies between technical competencies,
- interpret experience narratives more accurately,
- handle diverse resume writing styles,
- improve classification robustness across heterogeneous datasets.

Their inclusion ensures higher classification accuracy and stronger generalization performance when processing resumes from candidates with different linguistic and professional backgrounds.

### **Integration of Fairness Evaluation Metrics**

Fairness evaluation metrics were incorporated as essential components of the framework to ensure compliance with ethical AI deployment principles. Automated recruitment systems must be capable of identifying unintended discrimination patterns across demographic groups to maintain transparency and accountability.

The selected fairness metrics include:

- Disparate Impact Ratio (DIR)
- Demographic Parity Difference (DPD)
- Equalized Odds Difference (EOD)

These metrics enable quantitative monitoring of demographic bias across protected attributes including gender, age group, and name-origin categories. Their integration ensures that classification performance improvements do not come at the cost of fairness violations in hiring decisions.

### **Incorporation of Bias Mitigation Algorithms**

Bias mitigation algorithms were integrated into the prediction pipeline to reduce discrimination effects during classification while preserving predictive performance. These mitigation techniques improve demographic balance in model outputs by adjusting training data distributions and learning representations.

The mitigation framework includes:

- sample re-weighting strategies for balancing class representation,
- counterfactual data augmentation for reducing attribute sensitivity,
- adversarial debiasing for fairness-aware representation learning.

The integration of these techniques ensures that the screening system produces equitable predictions across demographic groups without significantly degrading classification accuracy.

### **Consideration of Ethical and Legal Constraints**

Since automated hiring systems operate within sensitive decision-making environments, ethical and regulatory compliance requirements played an important role during feature selection. The framework avoids the direct use of protected demographic attributes for decision-making while still allowing fairness evaluation through controlled analytical modules.

This design strategy ensures:

- transparency in automated screening decisions,
- compliance with equal opportunity hiring principles,
- reduction of unintended algorithmic discrimination,
- alignment with responsible AI deployment practices.

Such considerations strengthen the suitability of the proposed system for deployment within modern recruitment analytics environments.

### **Future-Oriented Deployment Components**



Certain advanced deployment-oriented components such as explainable AI dashboards and real-time Applicant Tracking System (ATS) integration were identified as future extensions due to additional infrastructure requirements and enterprise-level implementation complexity. Although these modules were not included in the current implementation stage, they represent important directions for extending the system toward production-ready recruitment platforms.

These future extensions may support:

- visualization of feature importance and classification confidence,
- recruiter-interactive fairness monitoring interfaces,
- automated resume ingestion through enterprise ATS platforms,
- large-scale cloud-based deployment pipelines,
- continuous fairness auditing during recruitment cycles.

The final feature configuration therefore represents a balanced hybrid architecture that combines interpretable classical machine learning techniques with high-performance transformer models and fairness-aware optimization strategies. This combination ensures that the system remains technically robust, ethically compliant, and scalable for future intelligent recruitment automation environments.

**Table 3.2: Feature Finalization after Constraint Analysis**

| Feature                     | Decision     | Reason                                              |
|-----------------------------|--------------|-----------------------------------------------------|
| TF-IDF Classification       | Selected     | Baseline interpretability and fast execution        |
| Transformer Classification  | Selected     | Context-aware semantic prediction                   |
| Fairness Metrics            | Selected     | Bias quantification required for ethical compliance |
| Bias Mitigation Algorithms  | Selected     | Improve demographic parity across groups            |
| Named Entity Recognition    | Selected     | Extract structured candidate information            |
| Model Evaluation Metrics    | Selected     | Ensure performance validation reliability           |
| Explainable AI Dashboard    | Future Scope | Requires visualization infrastructure support       |
| Real-Time ATS Integration   | Future Scope | Requires enterprise recruitment APIs                |
| Multilingual Resume Support | Future Scope | Needs extended language model training              |
| Cloud Deployment Pipeline   | Future Scope | Requires distributed computing resources            |

In addition to feature feasibility analysis, constraint-based filtering was applied to ensure that the selected architecture remains suitable for academic implementation while maintaining extensibility for industrial deployment scenarios. The primary constraints considered during feature selection are summarized below.

**Table 3.3: Constraint-Based Feature Selection Considerations**

| Constraint Type           | Impact on Feature Selection                                                                                            |
|---------------------------|------------------------------------------------------------------------------------------------------------------------|
| Technical Constraints     | Limited availability of large-scale labeled resume datasets restricted training of very deep transformer architectures |
| Computational Constraints | GPU memory requirements influenced transformer fine-tuning depth and batch size selection                              |
| Ethical Constraints       | Mandatory inclusion of fairness metrics and mitigation modules in prediction pipeline                                  |
| Legal Constraints         | Avoidance of direct demographic decision-making features ensured compliance with hiring fairness policies              |
| Data Privacy Constraints  | Use of synthetic dataset prevented exposure of sensitive personal recruitment information                              |
| Scalability Constraints   | Modular architecture enabled future integration with enterprise recruitment systems                                    |

The final feature configuration therefore represents a balanced hybrid framework that combines interpretable classical machine learning models with high-performance transformer architectures and fairness-aware optimization techniques. This ensures that the system remains suitable for academic experimentation while also supporting future deployment in real-world intelligent recruitment environments.

### 3.3 Model Training Strategy

Model training was performed using supervised classification techniques. The dataset was divided into training and testing subsets using an 80:20 split to ensure generalization capability

of the models.

Classical machine learning classifiers were trained using TF-IDF vector representations, while transformer models were fine-tuned using contextual token embeddings derived from pre-trained language models.

The optimization objective for classification models is defined as:

$$\theta^* = \arg \min_{\theta} \mathcal{L}(y, f(x; \theta))$$

where

- $x$  represents resume feature vectors
- $y$  represents class labels
- $f(x; \theta)$  represents classification function
- $\mathcal{L}$  represents cross-entropy loss

### 3.4 Transformer Fine-Tuning Configuration

Transformer-based models such as BERT and RoBERTa were fine-tuned using task-specific resume classification datasets. The models utilize contextual embeddings extracted from the [CLS] token representation.

The fine-tuning parameters used in the training process include:

- Learning rate:  $2 \times 10^{-5}$
- Batch size: 16
- Epochs: 3
- Optimizer: AdamW
- Maximum token length: 512

The classification probability is computed as:

$$P(y|x) = \text{Softmax}(Wh_{CLS} + b)$$

where  $h_{CLS}$  represents contextual embedding from transformer encoder.

### 3.5 Bias Detection Methodology

Bias detection is performed using fairness evaluation metrics across protected demographic attributes. The fairness evaluation module measures whether candidate selection probabilities differ significantly between privileged and unprivileged groups.

The fairness objective can be expressed as:

$$\Delta = |P(\hat{Y} = 1|A = a) - P(\hat{Y} = 1|A = b)|$$

where  $A$  represents demographic attribute categories.

### 3.6 Bias Mitigation Strategy

Bias mitigation techniques were implemented to reduce demographic disparity in classification outputs while preserving predictive accuracy.

The mitigation strategies applied include:

- Sample re-weighting
- Counterfactual data augmentation
- Adversarial debiasing

The adversarial debiasing objective function is defined as:

$$\min_{\theta_f} \max_{\theta_a} \mathcal{L}_{task} - \lambda \mathcal{L}_{adv}$$

where

- $\mathcal{L}_{task}$  represents classification loss
- $\mathcal{L}_{adv}$  represents adversarial fairness loss
- $\lambda$  controls fairness-performance tradeoff

### 3.7 Implementation Workflow

The implementation pipeline follows a structured processing sequence beginning from dataset preparation and ending with fairness-aware classification output.

**Table 3.4: Implementation Workflow of Resume Screening Framework**

| Stage                   | Description                                           |
|-------------------------|-------------------------------------------------------|
| Dataset Preparation     | Structured resume dataset with demographic attributes |
| Text Preprocessing      | Tokenization, stopword removal, normalization         |
| Feature Extraction      | TF-IDF vectors and transformer embeddings             |
| Model Training          | LR, SVM, RF, BERT, RoBERTa training                   |
| Performance Evaluation  | Accuracy, Precision, Recall, F1-score computation     |
| Bias Detection          | DIR, DPD, EOD fairness metric calculation             |
| Bias Mitigation         | Re-weighting and adversarial debiasing applied        |
| Final Prediction Output | Fair shortlist classification results generated       |

### 3.8 System Scalability Considerations

The proposed automated resume screening framework is designed using a modular and extensible architecture so that it can be adapted for large-scale recruitment environments without requiring structural redesign. Each component of the pipeline, including text preprocessing, feature extraction, classification, fairness evaluation, and bias mitigation, operates as an independent processing block. This modularity enables flexible upgrades, parallel execution, and distributed deployment in enterprise hiring systems.

To support large-volume resume processing, the architecture allows batch-based as well as stream-based document ingestion pipelines. This ensures that the system can efficiently handle recruitment datasets containing thousands or millions of candidate profiles while maintaining prediction accuracy and fairness compliance.

The framework is designed to support future extensions including:

- integration with Applicant Tracking Systems (ATS) for automated resume ingestion
- deployment as recruitment decision-support REST APIs

- explainable AI visualization dashboards for recruiter transparency
- real-time resume ranking pipelines for continuous candidate evaluation
- cloud-based deployment using distributed GPU-enabled environments
- multilingual resume processing using transformer-based language adaptation
- adaptive fairness monitoring modules for regulatory compliance auditing

The architecture also supports horizontal scalability through distributed model inference, allowing multiple classification requests to be processed simultaneously. Transformer-based components can be deployed using containerized environments such as Docker or Kubernetes, enabling scalable inference services across enterprise infrastructure.

In addition, fairness evaluation modules can be extended into continuous monitoring pipelines that automatically detect demographic imbalance trends during recruitment cycles. This enables organizations to maintain compliance with ethical AI standards and employment regulations while preserving transparency in automated hiring decisions.

The proposed framework can further be extended to support hybrid deployment modes including:

- on-premise secure recruitment analytics platforms
- cloud-hosted AI recruitment assistants
- edge-based lightweight screening models for fast preliminary filtering

These extensions ensure that the framework remains adaptable, scalable, and suitable for deployment in real-world intelligent recruitment automation environments.

## CHAPTER-4

### Results Analysis and Validation

#### 4.1 Model Performance

The performance evaluation of the proposed automated resume screening framework was carried out using standard classification metrics including accuracy, precision, recall, and F1-score. These evaluation metrics provide a comprehensive assessment of model prediction capability by measuring correctness, reliability, sensitivity, and balance between true positive and false positive classifications.

Accuracy represents the proportion of correctly classified resumes across the entire dataset. Precision measures how many selected candidates are actually relevant, while recall evaluates the ability of the model to identify all suitable candidates from the dataset. The F1-score provides a harmonic balance between precision and recall and is particularly useful when class distributions are uneven in recruitment datasets.

The comparative performance results obtained from classical machine learning models and transformer-based architectures are presented in Table 4.1.


**Table 4.1: Model Performance Comparison**

| Model        | Accuracy | Precision | Recall | F1-Score |
|--------------|----------|-----------|--------|----------|
| LR + TF-IDF  | 82.3%    | 0.81      | 0.82   | 0.81     |
| TF-IDF + RF  | 84.1%    | 0.83      | 0.84   | 0.83     |
| SVM + TF-IDF | 83.7%    | 0.82      | 0.84   | 0.83     |
| BERT         | 91.7%    | 0.92      | 0.91   | 0.92     |
| RoBERTa      | 93.2%    | 0.93      | 0.93   | 0.93     |

The results clearly indicate that transformer-based architectures outperform classical TF-

IDF-based models in resume classification tasks. Traditional machine learning approaches rely primarily on keyword-frequency representations and therefore cannot effectively capture contextual semantic relationships between skills, experience descriptions, and professional qualifications.

In contrast, transformer architectures such as BERT and RoBERTa utilize contextual embedding mechanisms that interpret word relationships within resume content. This improves classification robustness across diverse resume writing styles and increases prediction reliability in real-world recruitment environments.



**Figure 4.1: Accuracy Comparison of Resume Screening Models**

Figure 4.1 illustrates that contextual transformer-based models achieve significantly higher classification accuracy compared to classical machine learning approaches. This confirms the effectiveness of deep contextual representations for automated resume understanding and candidate screening tasks.

#### sectionBias Before Mitigation

Automated recruitment systems may unintentionally learn demographic patterns present in historical hiring datasets. Therefore, fairness evaluation was performed before applying mitigation techniques in order to quantify bias across multiple protected attributes including gender, race or ethnicity, age group, and name-origin indicators. Since resume screening decisions directly influence candidate shortlisting outcomes, identifying demographic imbalance at the baseline stage is essential for ensuring ethical reliability of the classification framework.

Bias measurements were computed using three standard fairness metrics: Disparate Impact Ratio (DIR), Demographic Parity Difference (DPD), and Equalized Odds Difference (EOD).



These indicators provide quantitative insight into whether classification decisions disproportionately favor or disadvantage specific candidate groups.

**Table 4.2: Fairness Metrics Before Mitigation**

| Bias Type        | DIR  | DPD  | EOD  |
|------------------|------|------|------|
| Gender           | 0.71 | 0.18 | 0.21 |
| Race / Ethnicity | 0.68 | 0.22 | 0.25 |
| Age              | 0.74 | 0.15 | 0.18 |
| Name-Based       | 0.72 | 0.17 | 0.20 |

The fairness evaluation results indicate the presence of measurable demographic bias in baseline classification outputs. In particular, race and ethnicity-based bias produced the lowest Disparate Impact Ratio values, suggesting unequal selection probability across candidate groups. These observations highlight the importance of incorporating fairness-aware optimization strategies within the resume screening pipeline.

Further analysis of the baseline fairness indicators provides the following important observations:

- The Disparate Impact Ratio values for all protected attributes remain below the recommended fairness threshold of 0.80, indicating statistically significant selection imbalance across demographic groups.
- The highest deviation was observed in race and ethnicity categories, which suggests that historical hiring patterns embedded within the dataset may influence classification outcomes more strongly for these attributes.
- Gender-based bias also shows noticeable disparity, indicating that the model may implicitly associate certain skill patterns or job roles with specific gender representations present in training data.
- Age-related bias appears comparatively lower than other attributes but still remains outside acceptable neutrality limits, confirming the presence of subtle age-related selection influence.
- Name-origin bias reflects linguistic or cultural correlations learned by the model from resume text features, which can indirectly affect classification fairness.

- The Demographic Parity Difference values indicate unequal probability of positive classification outcomes between protected and non-protected groups, which further confirms imbalance in prediction distributions.
- Equalized Odds Difference values demonstrate variation in true positive rates across demographic groups, suggesting that the classifier performs inconsistently when evaluating candidates belonging to different identity categories.

These findings confirm that baseline classification models capture hidden demographic signals from resume content even when protected attributes are not explicitly used as input features. Such behaviour is common in automated recruitment systems trained on historical datasets that reflect past organizational hiring preferences.

Therefore, the presence of measurable bias across multiple fairness indicators justifies the integration of bias mitigation algorithms in the subsequent stage of the screening framework. Addressing these disparities ensures that classification decisions remain aligned with ethical AI deployment principles and equal opportunity recruitment practices.

## 4.2 Bias Mitigation Results

To improve fairness performance, bias mitigation algorithms including adversarial debiasing, reweighting techniques, and counterfactual data augmentation were integrated into the classification workflow. These techniques reduce demographic sensitivity while preserving predictive accuracy.


The fairness improvements obtained after mitigation are summarized in Table 4.3.

**Table 4.3: Fairness Metrics After Mitigation**

| Bias Type        | Technique                   | DIR Improvement | DPD Improvement |
|------------------|-----------------------------|-----------------|-----------------|
| Gender           | Adversarial Debiasing       | 0.71 to 0.89    | 0.18 to 0.07    |
| Race / Ethnicity | Reweighting + CDA           | 0.68 to 0.87    | 0.22 to 0.09    |
| Age              | Counterfactual Augmentation | 0.74 to 0.91    | 0.15 to 0.06    |
| Name-Based       | Adversarial Debiasing       | 0.72 to 0.88    | 0.17 to 0.07    |

The results demonstrate a consistent improvement in fairness indicators across all evaluated demographic attributes. The Disparate Impact Ratio values moved closer to the fairness thresh-

old of 1.0, while Demographic Parity Difference values reduced significantly after mitigation. This confirms that fairness-aware optimization techniques successfully improved demographic neutrality without introducing instability in classification performance.



**Figure 4.2: DIR Before and After Bias Mitigation**

Figure 4.2 visually illustrates the improvement in fairness scores after applying mitigation strategies. The upward shift in DIR values across all demographic attributes confirms that the classification system achieved greater demographic balance following bias correction.

### 4.3 Validation

The proposed resume screening framework was validated using both classification performance metrics and fairness evaluation indicators. Classification metrics verified prediction accuracy and model reliability, while fairness metrics ensured that the system maintained demographic neutrality across protected candidate groups. This dual-layer validation strategy was necessary because automated hiring systems must satisfy both technical performance requirements and ethical compliance standards.

The validation procedure included the following evaluation components:


- Accuracy analysis to measure overall correctness of candidate classification.
- Precision evaluation to ensure selected candidates were relevant to job-role requirements.
- Recall measurement to confirm that suitable candidates were not incorrectly excluded.

- F1-score assessment to maintain balanced performance across uneven class distributions.
- Fairness metric verification using Disparate Impact Ratio (DIR), Demographic Parity Difference (DPD), and Equalized Odds Difference (EOD).
- Bias mitigation effectiveness testing to confirm demographic neutrality improvements.

The validation results confirm that transformer-based architectures significantly improve classification accuracy compared to traditional machine learning models. At the same time, fairness-aware optimization techniques successfully reduced bias levels without compromising predictive performance. This balance between accuracy and fairness demonstrates the suitability of the proposed framework for deployment in intelligent recruitment automation environments.

Additional validation observations include:

- Transformer embeddings improved semantic interpretation of candidate skillsets, resulting in higher classification confidence scores.
- Bias mitigation algorithms improved DIR values across all protected attributes, bringing predictions closer to fairness compliance thresholds.
- Classification stability remained consistent after mitigation, confirming that fairness optimization did not introduce model degradation.
- The hybrid architecture maintained interpretability through classical baseline models while benefiting from contextual understanding provided by transformer models.
- The framework supports scalable integration with enterprise Applicant Tracking Systems in future deployment scenarios.



**Figure 4.3: Validation Workflow of the Resume Screening Framework**

Figure 4.3 illustrates the structured validation pipeline adopted in the proposed system. The workflow confirms that classification performance evaluation and fairness auditing were integrated sequentially to ensure balanced optimization of predictive reliability and ethical compliance.

## 4.4 Interpretation


The experimental results indicate that contextual transformer models such as BERT and RoBERTa provide superior resume classification capability due to their ability to capture semantic relationships within candidate profiles. These architectures interpret contextual dependencies between skills, qualifications, certifications, and experience descriptions more effectively than traditional keyword-based methods.

However, transformer-based systems may also learn hidden demographic correlations present in training datasets, which makes fairness evaluation an essential component of the screening pipeline. Without fairness monitoring, such correlations could lead to unintended bias in automated recruitment recommendations.

Bias mitigation algorithms improved demographic parity across candidate groups while preserving acceptable classification accuracy levels. This confirms that fairness-aware machine learning techniques can be effectively integrated into automated hiring systems without degrading predictive reliability.

The interpretation of experimental findings leads to the following key insights:

- Contextual embedding models significantly improve classification robustness across heterogeneous resume formats.
- Classical machine learning models provide transparent baseline interpretability useful for recruiter-level validation.
- Fairness-aware optimization reduces demographic imbalance without affecting semantic classification quality.
- Hybrid architectures improve trustworthiness of automated screening decisions in real-world hiring environments.
- The framework demonstrates compatibility with responsible AI deployment principles required for modern recruitment analytics platforms.



**Figure 4.4: Hybrid Interpretation Strategy for Resume Classification**

Figure 4.4 illustrates how the hybrid architecture combines interpretable classical machine learning models with transformer-based contextual representations to produce accurate and fairness-aware classification outcomes.

Overall, the results validate that the proposed hybrid architecture combining classical machine learning models, transformer-based embeddings, and fairness-aware optimization strategies provides an accurate, interpretable, and ethically reliable solution for automated resume

screening applications. The framework therefore represents a practical step toward trustworthy AI-assisted recruitment systems capable of supporting scalable and unbiased hiring decisions.

## **CHAPTER-5**

### **Conclusion and Future Work**

#### **5.1 Conclusion**

This project presented the design and implementation of an automated resume screening framework using Natural Language Processing techniques with a strong focus on fairness evaluation and bias mitigation. The proposed system integrates classical machine learning models, transformer-based contextual embeddings, and fairness-aware optimization techniques to improve recruitment decision support while maintaining demographic neutrality across candidate groups.

The primary objective of this research was to investigate whether automated resume classification systems can achieve both high predictive accuracy and ethical fairness simultaneously. Experimental evaluation confirmed that transformer-based architectures such as BERT and RoBERTa significantly improved classification performance compared to traditional TF-IDF-based machine learning models.

However, fairness analysis revealed that higher classification accuracy alone does not guarantee unbiased decision-making. The results demonstrated measurable bias across gender, race or ethnicity, age group, and name-origin attributes when baseline models were evaluated using fairness metrics.

The major outcomes of the project are summarized as follows:

- A structured NLP-based resume preprocessing pipeline was developed to convert unstructured resume text into machine-readable representations.
- Classical machine learning classifiers including Logistic Regression, Support Vector Machine, and Random Forest were implemented as interpretable baseline models.



- Transformer-based architectures including BERT and RoBERTa improved semantic classification accuracy across heterogeneous resume formats.
- Fairness evaluation metrics including Disparate Impact Ratio (DIR), Demographic Parity Difference (DPD), and Equalized Odds Difference (EOD) were integrated into the evaluation pipeline.
- Bias mitigation algorithms such as sample re-weighting, adversarial debiasing, and counterfactual data augmentation successfully improved fairness scores across demographic groups.
- A hybrid architecture combining interpretability and contextual learning capability was established for responsible automated hiring support systems.

Overall, the results confirm that fairness-aware resume screening frameworks can achieve reliable classification performance while reducing unintended demographic bias in automated recruitment pipelines.

## **5.2 Validation of Project Objectives**

The objectives defined at the beginning of the project were successfully achieved through systematic implementation and evaluation of the proposed framework.

**Table 5.1: Validation of Project Objectives**

| Project Objective                                 | Validation Outcome                                             |
|---------------------------------------------------|----------------------------------------------------------------|
| Resume dataset preprocessing                      | Achieved through NLP-based cleaning and normalization pipeline |
| Feature extraction using TF-IDF and transformers  | Achieved through hybrid embedding architecture                 |
| Implementation of classification models           | Achieved using LR, SVM, RF, BERT, and RoBERTa                  |
| Fairness evaluation across demographic attributes | Achieved using DIR, DPD, and EOD metrics                       |
| Bias mitigation integration                       | Achieved using adversarial debiasing and reweighting methods   |
| Development of ethical screening framework        | Achieved through fairness-aware optimization pipeline          |

The validation results confirm that both technical and ethical objectives of the project were successfully accomplished.

### 5.3 Research Contributions of the Project

This work contributes to the development of responsible AI-based recruitment systems by integrating contextual language understanding with fairness-aware optimization techniques.

The major research contributions are summarized below:

- Development of a hybrid resume classification framework combining classical machine learning and transformer architectures.
- Integration of fairness evaluation metrics directly within the recruitment prediction pipeline.
- Implementation of multiple bias mitigation strategies to improve demographic neutrality of classification outputs.
- Comparative performance analysis of interpretable and contextual models for resume screening tasks.

- Demonstration of accuracy–fairness trade-off behaviour in automated hiring systems.
- Design of a scalable architecture suitable for future Applicant Tracking System integration.

These contributions support the advancement of trustworthy automated recruitment analytics systems.

## 5.4 Practical Significance of the Proposed Framework

The proposed framework provides several practical advantages for intelligent recruitment automation environments.

First, the hybrid architecture enables organizations to improve candidate shortlisting accuracy without sacrificing transparency.

Second, fairness evaluation ensures compliance with ethical AI deployment standards and equal opportunity hiring regulations.

Third, bias mitigation techniques improve trustworthiness of automated screening recommendations.

Fourth, transformer-based embeddings allow the system to handle diverse resume formats and writing styles effectively.

Finally, the modular design supports integration with enterprise Applicant Tracking Systems and HR analytics platforms.

## 5.5 Ethical Impact and Responsible AI Contribution

Automated recruitment systems must ensure fairness, transparency, and accountability to prevent unintended discrimination against protected demographic groups.

The proposed framework contributes to responsible AI deployment by:

- identifying demographic bias before deployment,
- quantifying fairness using standard evaluation metrics,
- reducing bias through mitigation algorithms,
- improving transparency in automated decision-making pipelines,

- supporting equal opportunity hiring practices.

These contributions strengthen the ethical reliability of AI-assisted hiring systems.

## **5.6 Deviation from Expected Results**

It was expected that transformer models would provide better classification accuracy compared to classical machine learning models. This expectation was successfully achieved, as RoBERTa produced the highest classification performance among all evaluated architectures.

However, it was also observed that higher classification accuracy did not automatically result in improved fairness. Transformer-based models demonstrated greater sensitivity to demographic proxy signals present in training datasets.

These findings confirm that contextual embeddings capture subtle correlations between candidate attributes and historical hiring patterns, which may introduce unintended bias if fairness evaluation is not incorporated during deployment.

Therefore, fairness-aware optimization remains an essential requirement for responsible implementation of automated recruitment systems.

## **5.7 Limitations of the Study**

Although the proposed framework produced reliable classification and fairness evaluation results, several limitations were identified during implementation.

- The dataset used for experimentation was limited in size compared to large-scale enterprise recruitment datasets.
- Real-time Applicant Tracking System integration was not implemented.
- Explainability dashboards were not fully developed.
- Multilingual resume processing was outside the scope of the current system.
- Intersectional bias analysis across multiple protected attributes was not performed.

These limitations provide direction for future research extensions.

## **5.8 Future Work**

The framework developed in this project provides a strong foundation for several future research and deployment improvements.

### **5.8.1 Testing on Real-World Resume Datasets**

Future work should evaluate the proposed framework using large-scale real-world recruitment datasets collected from enterprise hiring platforms.

### **5.8.2 Explainable AI Dashboard Development**

Interactive visualization dashboards can be developed to display feature importance scores and fairness indicators for HR decision-makers.

### **5.8.3 Intersectional Bias Analysis**

Future studies may evaluate combined demographic attributes such as gender-age and ethnicity-location relationships to detect intersectional bias patterns.

### **5.8.4 Web-Based Resume Screening Platform**

A web-based deployment interface can be developed to allow recruiters to upload resumes and receive automated classification predictions with fairness reports.

### **5.8.5 Multilingual Resume Processing**

Future versions of the system may incorporate multilingual transformer models to support resume classification across global recruitment environments.

### **5.8.6 Human-in-the-Loop Decision Support**

Integration of human validation stages can improve reliability for borderline classification cases and reduce automated decision risks.

## **5.9 Scalability of the Proposed Framework**

The modular architecture allows deployment across:

- small-scale academic recruitment datasets,
- enterprise-level Applicant Tracking Systems,
- cloud-based HR analytics platforms,
- cross-organizational hiring environments.

This scalability ensures long-term applicability of the proposed system.

## **5.10 Final Summary**

This chapter presented the conclusions drawn from the implementation and evaluation of the automated resume screening framework. The experimental results demonstrate that transformer-based contextual embeddings significantly improve classification accuracy, while fairness-aware optimization techniques reduce demographic bias in automated hiring decisions.

The proposed hybrid architecture therefore provides a practical foundation for developing accurate, interpretable, and ethically responsible AI-assisted recruitment systems capable of supporting future intelligent hiring platforms.

## REFERENCES

1. Dastin, J. (2018). Amazon scraps secret AI recruiting tool that showed bias against women. Reuters.
2. Raghavan, M., Barocas, S., Kleinberg, J., and Levy, K. (2020). Mitigating bias in algorithmic hiring: Evaluating claims and practices. ACM FAccT.
3. Maheshwary, S. and Misra, S. (2018). Matching resumes to jobs via deep learning. WWW.
4. Bolukbasi, T., Chang, K., Zou, J., Saligrama, V., and Kalai, A. (2016). Man is to computer programmer as woman is to homemaker? Debiasing word embeddings. NeurIPS.
5. Caliskan, A., Bryson, J., and Narayanan, A. (2017). Semantics derived automatically from language corpora contain human-like biases. Science.
6. Dwork, C., Hardt, M., Pitassi, T., Reingold, O., and Zemel, R. (2012). Fairness through awareness. ITCS.
7. Hardt, M., Price, E., and Srebro, N. (2016). Equality of opportunity in supervised learning. NeurIPS.
8. Chouldechova, A. (2017). Fair prediction with disparate impact. Big Data.
9. Feldman, M., Friedler, S., Moeller, J., Scheidegger, C., and Venkatasubramanian, S. (2015). Certifying and removing disparate impact. KDD.
10. Kamiran, F. and Calders, T. (2012). Data preprocessing techniques for classification without discrimination. Knowledge and Information Systems.
11. Zhang, B., Lemoine, B., and Mitchell, M. (2018). Mitigating unwanted biases with adversarial learning. AAAI/ACM AIES.
12. Zafar, M., Valera, I., Rodriguez, M., and Gummadi, K. (2017). Fairness beyond disparate treatment and disparate impact. WWW.
13. Devlin, J., Chang, M., Lee, K., and Toutanova, K. (2019). BERT: Pre-training of deep bidirectional transformers for language understanding. NAACL-HLT.
14. Liu, Y. et al. (2019). RoBERTa: A robustly optimized BERT pretraining approach.
15. Bender, E., Gebru, T., McMillan-Major, A., and Shmitchell, S. (2021). On the dangers of stochastic parrots. ACM FAccT.

# CHAPTER-A

## Appendix

### A.1 Sample Resume Fields

A resume in the dataset contains the following fields:

- Candidate Name
- Professional Summary
- Skills
- Work Experience
- Education
- Projects
- Certifications
- Shortlisting Label

These structured attributes were extracted from unstructured resume documents using Natural Language Processing preprocessing techniques. Each field plays an important role in enabling accurate classification of candidate suitability for target job roles.

The professional summary section provides contextual background information about the candidate's expertise and career orientation. The skills section serves as the primary feature source for technical keyword extraction and competency analysis. Work experience contributes temporal and domain relevance indicators, while education provides qualification-level signals required for screening decisions.

Projects and certifications further strengthen candidate evaluation by providing evidence of applied technical capability and specialization areas. Finally, the shortlisting label represents the classification target variable used for training machine learning and transformer-based prediction models.

### A.2 Dataset Preparation Workflow

Before training the classification models, resume data was processed through a structured preprocessing pipeline. This pipeline ensured consistency across documents and improved feature extraction accuracy.

The preprocessing workflow included:



- Removal of special characters and formatting inconsistencies
- Conversion of text into lowercase format for normalization
- Tokenization of resume content into meaningful word units
- Stop-word removal to eliminate irrelevant linguistic elements
- Lemmatization for reducing words to base forms
- Feature vector generation using TF-IDF and transformer embeddings

This preprocessing stage significantly improved classification performance and reduced noise present in raw resume datasets.

## **A.3 User Manual**

### **A.3.1 System Requirements**

The minimum hardware and software configuration required to execute the proposed resume screening framework is listed below:

- Operating System: Windows / Linux / macOS
- Programming Language: Python 3.8 or higher
- Required Libraries:
  - pandas
  - numpy
  - scikit-learn
  - transformers
  - torch
  - spaCy
  - matplotlib
  - seaborn
- Development Environment: Jupyter Notebook / VS Code / PyCharm

These requirements ensure compatibility with both classical machine learning pipelines and transformer-based contextual embedding architectures.

### **A.3.2 Steps to Run the Project**

1. Install Python.
2. Install required libraries such as pandas, scikit-learn, transformers, torch, and spaCy.
3. Load the resume dataset.
4. Run preprocessing.
5. Train machine learning models.
6. Evaluate classification results.
7. Calculate fairness metrics.
8. Apply bias mitigation techniques.
9. Compare results before and after mitigation.

After completing these steps, the system generates classification performance metrics and fairness evaluation reports that support responsible recruitment decision-making.

### **A.3.3 Execution Workflow of the Framework**

The operational workflow of the resume screening framework follows a sequential pipeline structure as described below:

1. Resume dataset ingestion
2. Text preprocessing and normalization
3. Feature extraction using TF-IDF and transformer embeddings
4. Model training using classical and deep learning classifiers
5. Prediction of candidate suitability labels
6. Fairness evaluation across demographic attributes
7. Bias mitigation using optimization techniques
8. Comparative analysis of results before and after mitigation

This workflow ensures both predictive accuracy and demographic neutrality in automated screening decisions.

### **A.3.4 System Outputs**

The proposed system generates the following outputs:

- Classification accuracy comparison charts
- Precision, recall, and F1-score evaluation tables
- Fairness metric evaluation reports
- Bias comparison graphs before mitigation
- Bias comparison graphs after mitigation
- Model performance visualizations
- Final recommendation-ready classification summaries

These outputs assist researchers and HR professionals in understanding model behaviour and validating fairness-aware screening performance.

## **A.4 Fairness Evaluation Module**

The fairness evaluation module measures demographic neutrality across protected attributes including gender, age group, race or ethnicity, and name-origin categories. The following fairness indicators are computed:

- Disparate Impact Ratio (DIR)
- Demographic Parity Difference (DPD)
- Equalized Odds Difference (EOD)

These metrics help detect unintended bias patterns in classification outputs and support ethical deployment of automated hiring systems.

## **A.5 Bias Mitigation Module**

Bias mitigation techniques were integrated into the classification pipeline to reduce demographic imbalance without affecting prediction reliability.

The implemented mitigation techniques include:

- Sample re-weighting for balanced dataset representation
- Counterfactual data augmentation for attribute neutrality
- Adversarial debiasing for fairness-aware feature learning

These techniques improve trustworthiness of automated screening recommendations.

## A.6 Achievements

- Developed a complete NLP-based resume screening framework.
- Compared classical ML and transformer-based models.
- Evaluated fairness across multiple demographic dimensions.
- Applied bias mitigation techniques.
- Prepared a structured project report according to university format.

In addition to the above contributions, the project successfully demonstrated the feasibility of integrating fairness-aware machine learning techniques into automated recruitment pipelines. The hybrid architecture combining interpretable classical models and contextual transformer embeddings improves both prediction accuracy and ethical reliability of screening decisions.

The implementation also provides a scalable foundation for future deployment in enterprise-level Applicant Tracking Systems and intelligent recruitment analytics platforms. Furthermore, the framework supports extension toward explainable AI interfaces that allow recruiters to understand model reasoning behind candidate selection outcomes.

Overall, the achievements of this project contribute toward the development of responsible AI-assisted hiring systems capable of improving recruitment efficiency while maintaining transparency and fairness across demographic groups.